

Cost-aware Stopping for Bayesian Optimization


Qian Xie,^{*1} Linda Cai,^{*2} Alexander Terenin,¹ Peter Frazier,¹ and Ziv Scully¹
¹Cornell University ²University of California, Berkeley



Abstract

In automated machine learning, scientific discovery, and other applications of Bayesian optimization, deciding when to stop evaluating expensive black-box functions is an important practical consideration. While several adaptive stopping rules have been proposed, in the cost-aware setting they lack guarantees ensuring they stop before incurring excessive function evaluation costs. We propose a *cost-aware stopping rule* for Bayesian optimization that adapts to varying evaluation costs and is free of heuristic tuning. Our rule is grounded in a theoretical connection to state-of-the-art cost-aware acquisition functions, namely the Pandora's Box Gittins Index (PBGI) and log expected improvement per cost. We prove a theoretical guarantee bounding the expected cumulative evaluation cost incurred by our stopping rule when paired with these two acquisition functions. In experiments on synthetic and empirical tasks, including hyperparameter optimization and neural architecture size search, we show that combining our stopping rule with the PBGI acquisition function consistently matches or outperforms other acquisition-function-stopping-rule pairs in terms of *cost-adjusted simple regret*, a metric capturing trade-offs between solution quality and cumulative evaluation cost.

Cost-aware Bayesian Optimization



Cost-adjusted simple regret:
$$\underbrace{\min_{1 \leq t \leq \tau} f(x_t) - \inf_{x \in X} f(x)}_{\text{simple regret}} + \underbrace{\sum_{t=1}^{\tau} c(x_t)}_{\text{cumulative cost}}$$

Goal: *Adaptively* select evaluations x_1, x_2, \dots and stop at time τ to minimize the expected cost-adjusted simple regret.

Existing Adaptive Stopping Rules

Simple heuristics: stop when the best observed value remains unchanged or improvement is not statistically significant.

Acquisition-based: stop when PI, EI or KG falls below a threshold.

Regret-based: stop when regret bounds drop below a threshold (with some probability) such as in UCB-LCB.

PBGI/LogEIPC Stopping Rule

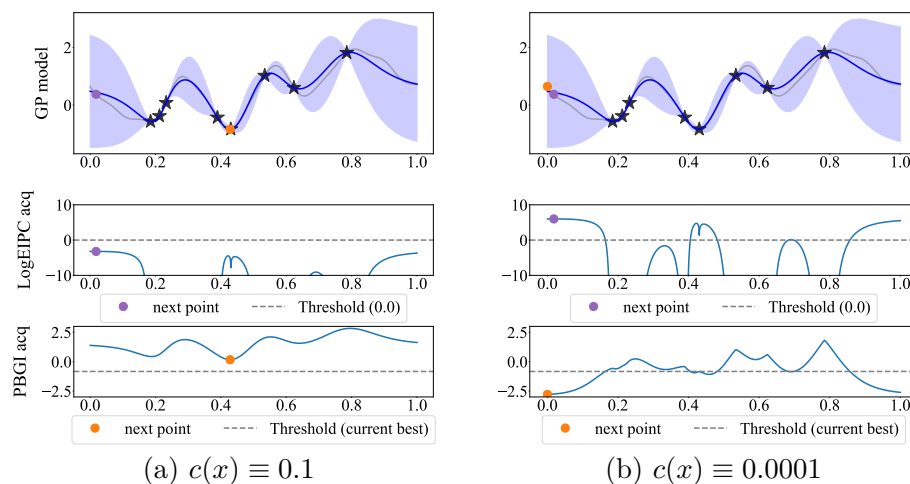
EI stopping rule. Stop when the expected improvement is no longer worth the unit cost: $\alpha_t^{\text{EI}}(x; y_{1:t}) \leq c$.

PBGI/LogEIPC stopping rule (this work). Stop when the Gittins index at *every* unevaluated point is at least the current best observed value:

$$\min_{x \in X \setminus \{x_1, \dots, x_t\}} \alpha_t^{\text{PBGI}}(x) \geq y_{1:t}^* \Leftrightarrow \max_{x \in X \setminus \{x_1, \dots, x_t\}} \alpha_t^{\text{LogEIPC}}(x; y_{1:t}) \leq 0.$$

Result (Weitzman, 1979). Under the *independent-value* setting, it is *Bayesian-optimal* when paired with the PBGI acquisition function.

Behavior Illustration



Theoretical Guarantee

Theorem 1 (No worse than stopping-immediately)

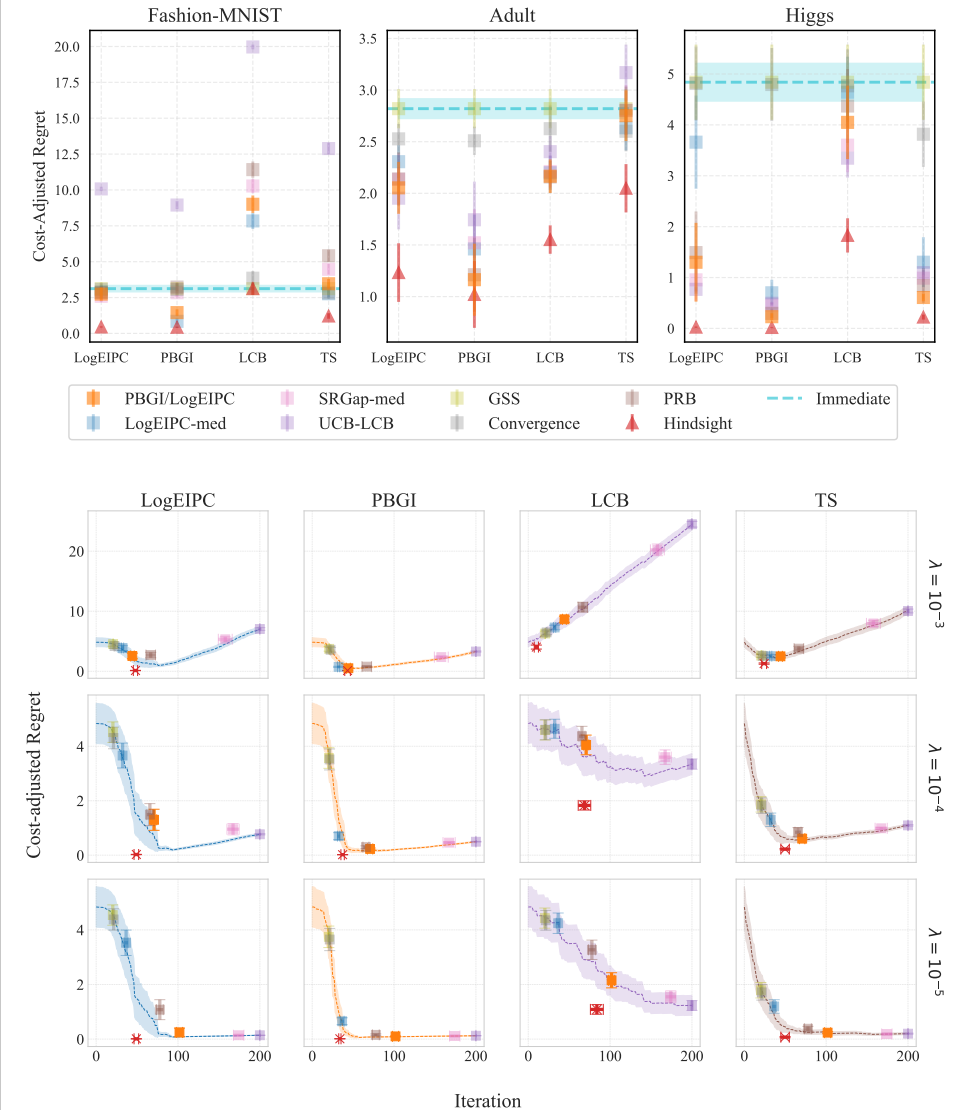
When optimizing a random function f with constant prior mean, using PBGI or LogEIPC with our stopping rule achieves cost-adjusted regret no worse than stopping immediately after the initial evaluation.

$$\mathbb{E} \left[y_{1:\tau}^* - \min_{x \in X} f(x) + \sum_{t=1}^{\tau} c(x_t) \right] \leq \mathbb{E} \left[y_1 - \min_{x \in X} f(x) + c(x_1) \right].$$

Key proof idea: Using our stopping rule, both PBGI and LogEIPC are guaranteed to evaluate only points whose one-step expected improvement is worth the evaluation cost before stopping.

Implication: Matches the best we can hope for in the worst case, and avoids over-spending — properties many cost-unaware rules lack.

Performance



Computation Time

