Cost-aware Bayesian Optimization via the Pandora's Box Gittins Indices

Qian Xie[1], Raul Astudillo[2], Peter Frazier[1], Ziv Scully[1], Alexander Terenin[1]

[1] Cornell University, [2] California Institute of Technology

Introduction to Bayesian optimization

Goal: optimize expensive-to-evaluate black-box function ∈ decision-making under uncertainty

An unknown random function $f: \mathcal{X} \to \mathbb{R}$ drawn from a Gaussian process prior

Gaussian process: infinite-dimensional generalization of multivariate normal distributions

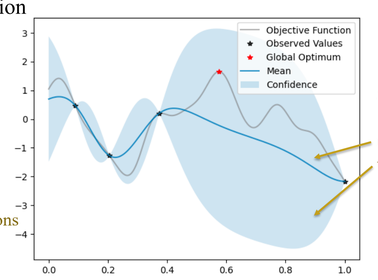

mean: prediction
variance: confidence/uncertainty

Trade-off between
- exploitation (high mean) and
- exploration (high uncertainty)

Applications:
Hyperparameter tuning
Drug discovery
Control design

x: hyperparameter/configuration

Objective: find global optimum $x^* = \operatorname{argmax}_{x \in \mathcal{X}} f(x)$

⇓

Objective: optimize best observed value at time T

$$\max_{\text{policy}} \mathbb{E} \max_{t=1,2,\dots,T} f(x_t)$$

Decision: evaluate a set of points

⇓

Decision: adaptively evaluate $x_1, x_2, \dots, x_T \in \mathcal{X}$ given time budget T

Why is Bayesian optimization hard?

Hard budget constraint

~~$t = 1$~~
~~$t = 2$~~
~~$t = 3$~~
~~$t = 4$~~
⋮
$t = T$

Correlated values

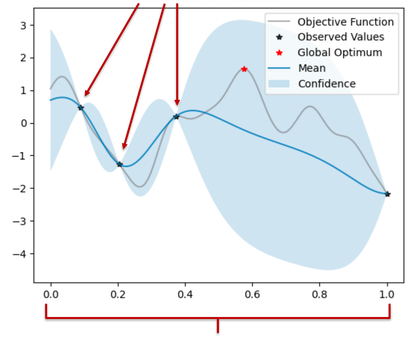

Continuous search domain

Evaluation costs handling

cheap	expensive
risk-seeking	risk-averse
exploration	exploitation
uniform	heterogeneous

Optimal policy unknown!

Connection with Pandora's box

special case of Markovian/Bayesian MAB

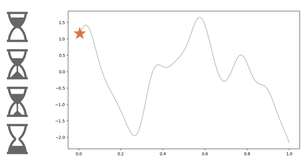

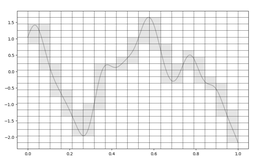

Continuous ⟹ Discrete

Correlated ⟹ Independent

Lagrangian relaxation

Hard budget constraint ⟹ Cost per sample

extension of [Aminian et al.'24]

Is Gittins index good? ⟸ Optimal policy: Gittins index

How to translate? [Weitzman'79]

Objective: maximize net utility

$$\max_{\text{policy}} \mathbb{E}\left(\max_{t=1,2,\dots,T} f(x_t) - \sum_{t=1}^{T} c(x_t)\right)$$

T: random stopping time

One-step heuristics: EI vs PBGI

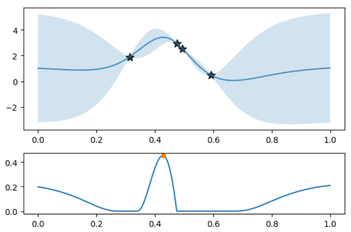

Expected improvement (EI)

popular

$\mathrm{EI}_{f|D}(x; y) = \mathbb{E}[((f|D)(x) - y)^+]$

EI policy: evaluate $\operatorname{argmax}_x \mathrm{EI}_{f|D}(x; y_{\text{best}})$

D: observed data, y_{best}: current best observed value

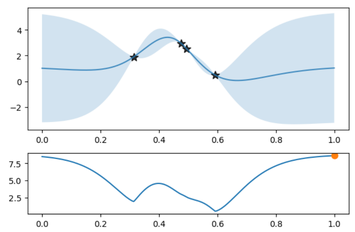

Pandora's box Gittins index (PBGI)

new

$g(x)$: solution to $\mathrm{EI}_{f|D}(x; g(x)) = \lambda$

PBGI policy: evaluate $\operatorname{argmax}_x g(x)$

λ: cost-per-sample (Lagrange multiplier)

Other acquisition functions:

simple
- Upper Confidence Bound (UCB)
- Thompson Sampling (TS)

unreliable
- Predictive Entropy Search
- Knowledge Gradient (KG)
- Multi-step Lookahead EI (MSEI)

slow

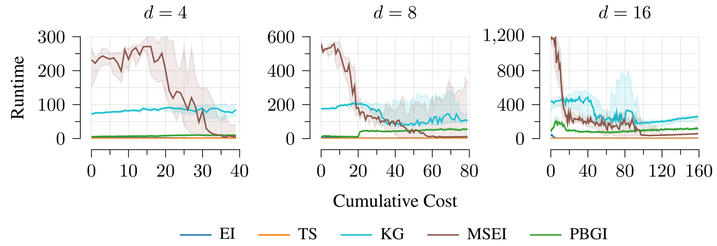

PBGI is easy to compute using bisection method!

Contour Plot

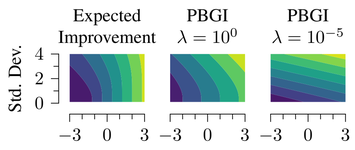

Connection with UCB?

Impact of λ

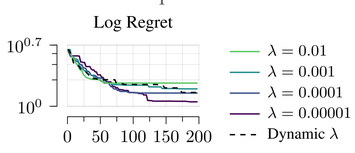

Smaller λ, higher exploration

Extension to heterogeneous costs

- Given cost function $c: \mathcal{X} \to \mathbb{R}^+$ and budget B
- Replace λ with $\lambda c(x)$ to compute $g(x)$ as PBGI

Baselines: arbitrarily bad
- EI Per Unit Cost (EIPC)
- Budgeted MSEI (BMSEI) slow

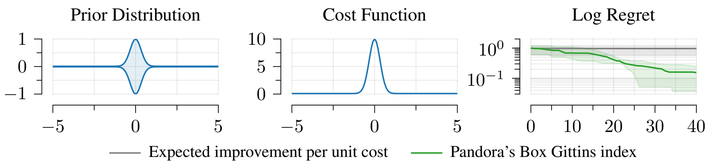

Experiment: Bayesian regret

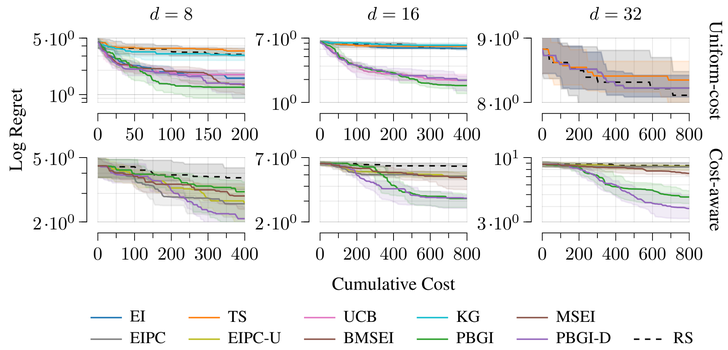

Experiment: synthetic benchmark

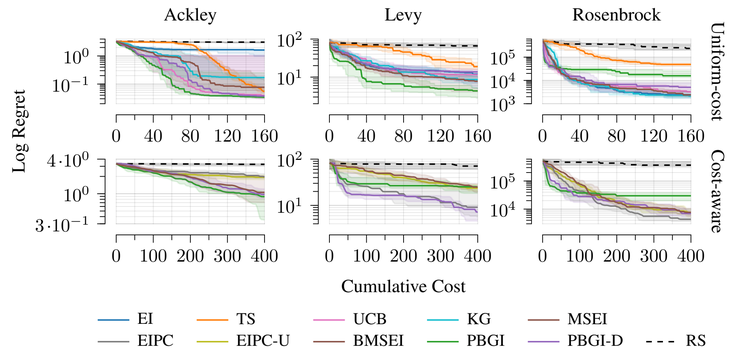

Experiment: empirical

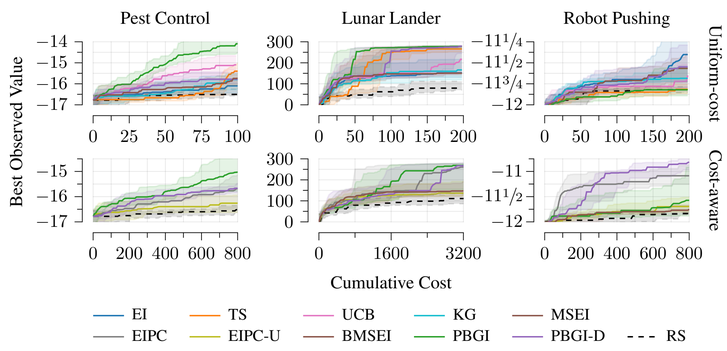

Future work

Extension to complex BO (freeze-thaw, multi-fidelity, function network, etc.) via Gittins variants ("golf" Markovian MAB, optional inspection, etc.)